




Convolutional Neural Networks and Applications in the NEXT Experiment

G. Díaz, J.A. Hernando, J. Renner
*IFGAE / Universidade de Santiago
de Compostela*

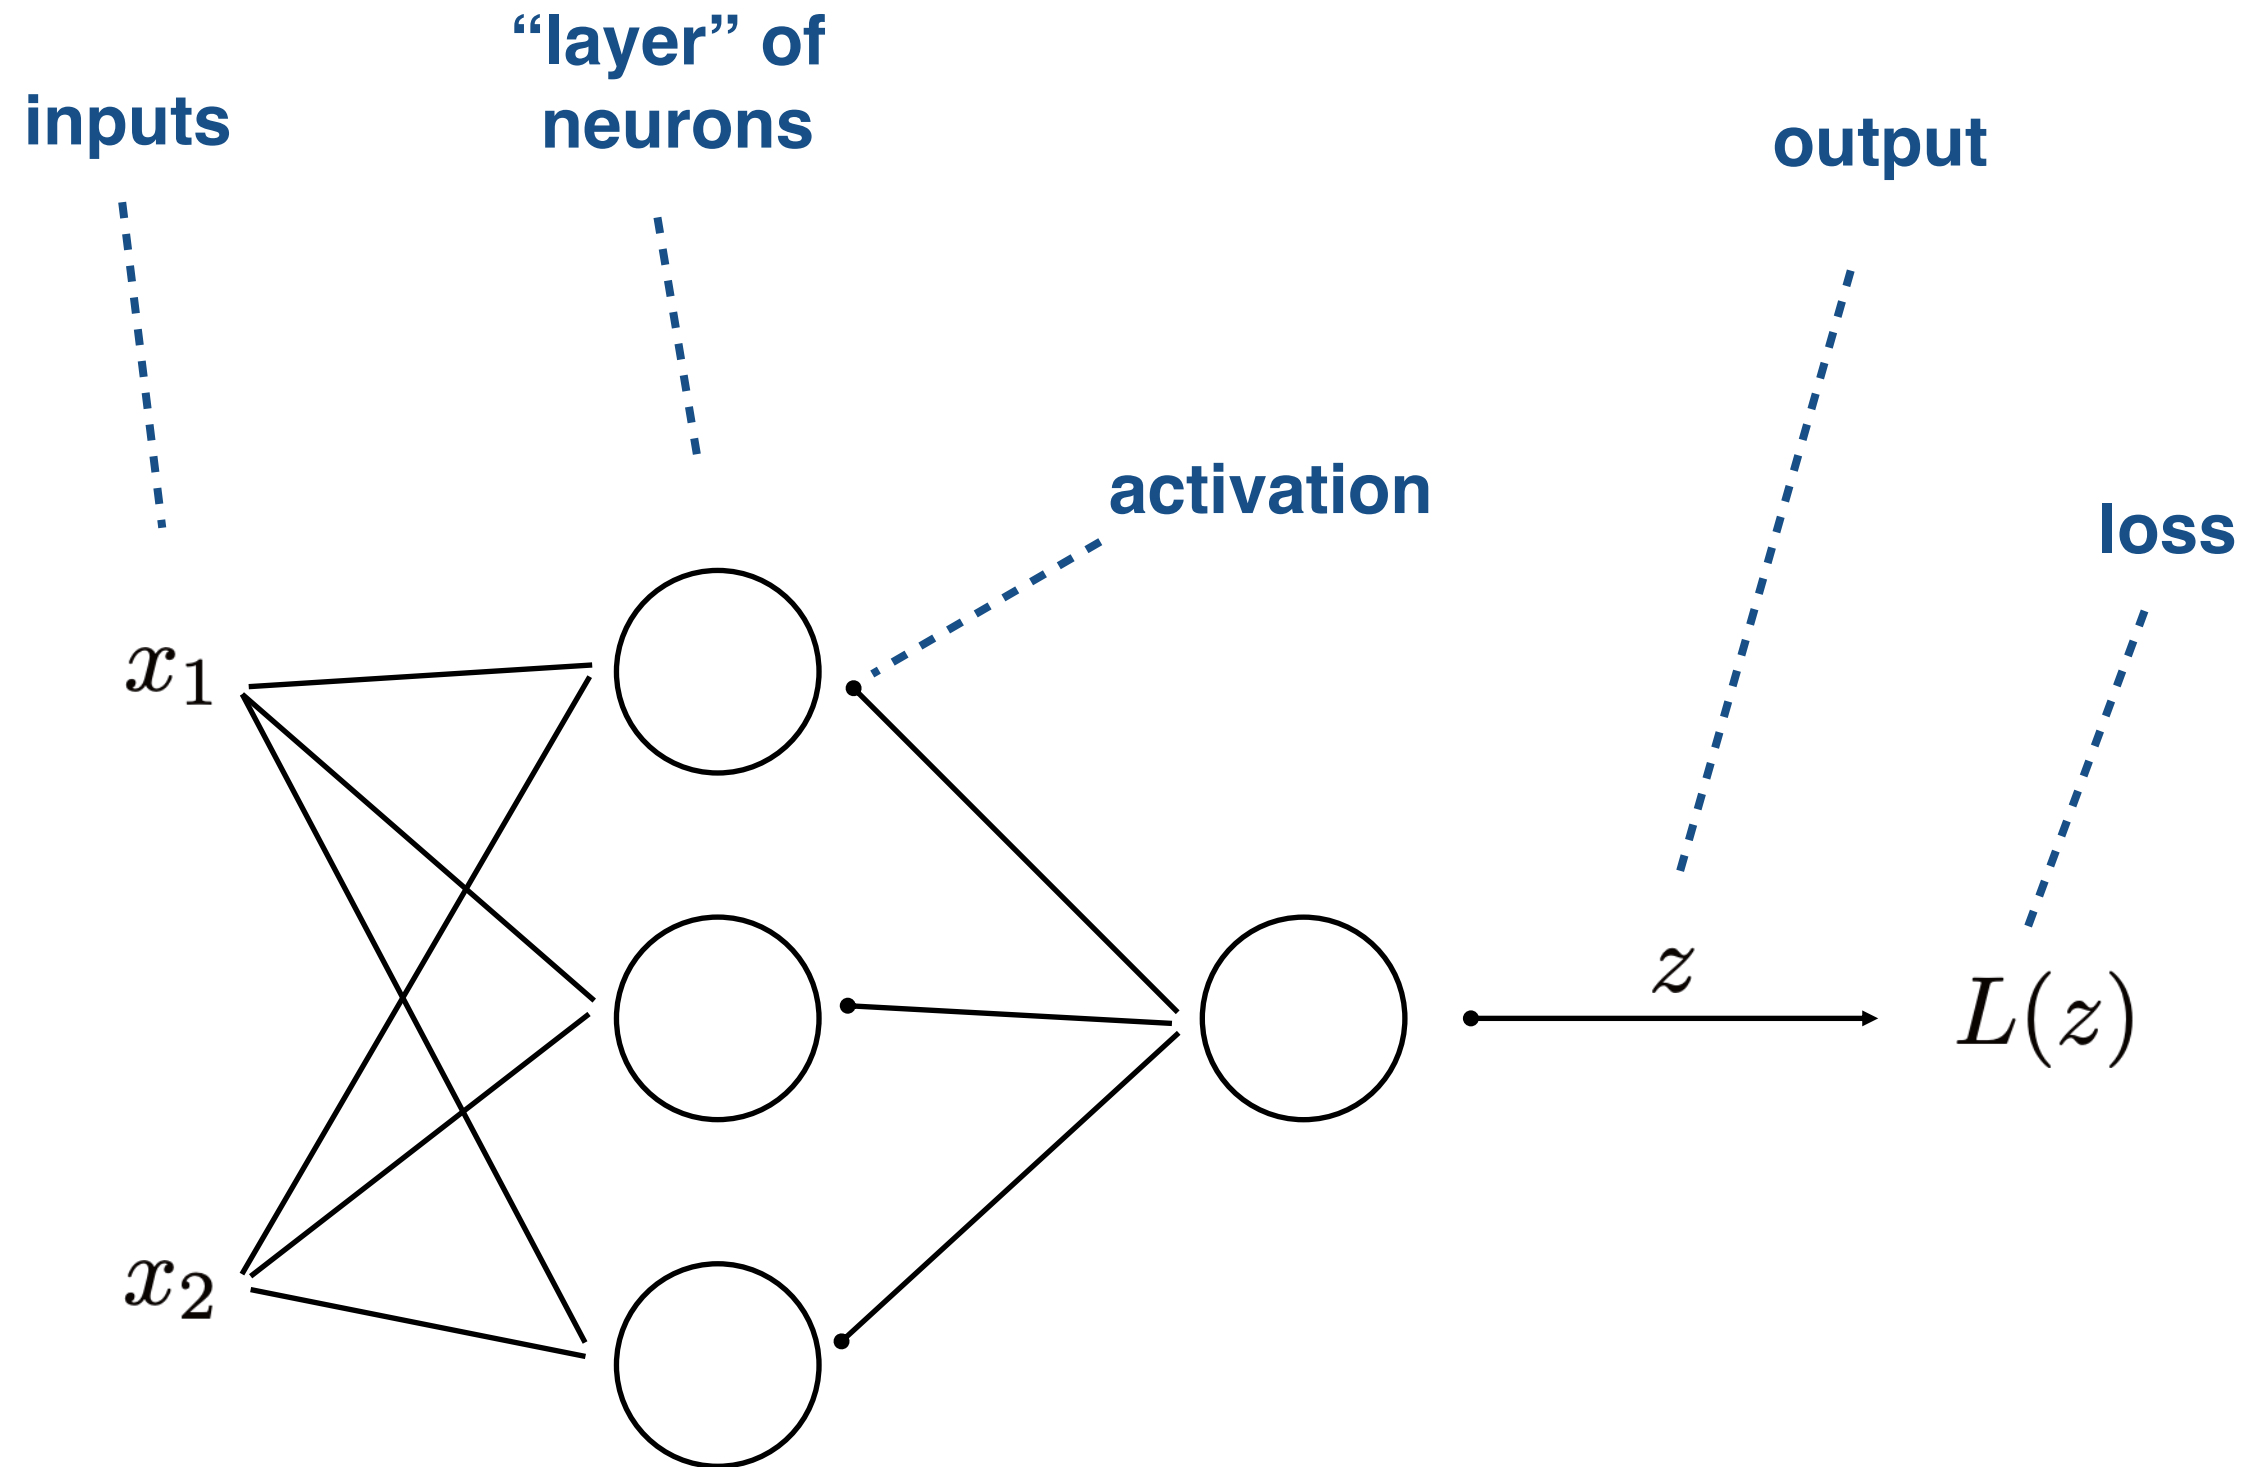
**Advanced Computing and
Machine Learning**
June 1, 2020

Review: last week

Neural networks

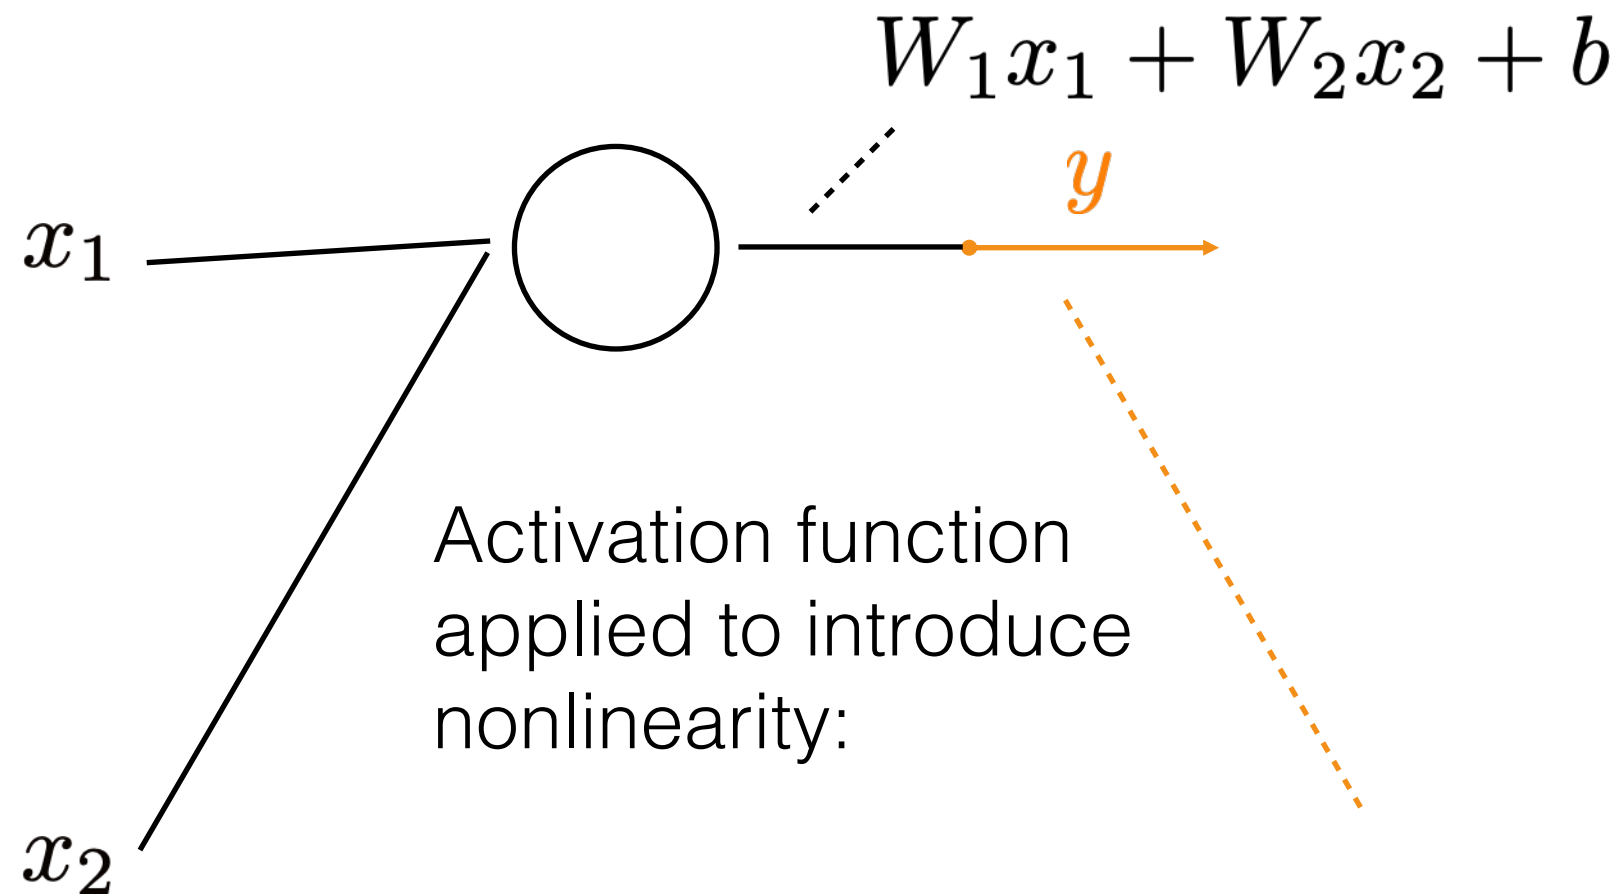


Neural networks



Neural networks

Each neuron contains a **weight** corresponding to each input, and a **bias**:




$$y = \sigma(W_1x_1 + W_2x_2 + b) = \frac{1}{1 + e^{-(W_1x_1 + W_2x_2 + b)}}$$

Neural networks


To minimize L , we want to push the entire network (all parameters) in the direction of $-\text{gradient}(L)$:

$$W_1 \rightarrow W_1 - \left(\frac{\partial L}{\partial W_1} \right) \epsilon$$



Neural networks

$$\left(\frac{\partial L}{\partial W_1}\right)\epsilon = \left(\frac{\partial y}{\partial W_1}\right) \left(\frac{\partial z}{\partial y}\right) \left(\frac{\partial L}{\partial z}\right)\epsilon$$




Note: $\partial z / \partial y$ will have to account for the functional forms of: **this neuron**


and this activation

Neural networks


- Note that backpropagation on graphs can be performed on a node-by-node basis:



Neural networks



Neural networks




- Given gradients dy_1 and dy_2
- Given f differentiable
- Compute gradients dW for parameter updates, and gradients dx_1 and dx_2 to continue backpropagation

$$dx_i = \sum_k \frac{\partial y_k}{\partial x_i} dy_k$$

Libraries such as PyTorch handle the mechanics of this automatically

Neural networks




- So far, we've considered:

$$f(\vec{x}, W) = W\vec{x} + \vec{b}$$


Let's try something else!

Convolutional Neural Networks (CNNs)

CNNs



Inputs, arranged in
a 2D matrix




KxL kernel of weights
(in this case 2x2)


CNNs

Multiply and add:


$$z_{11} = W_{11}x_{11} + W_{12}x_{12} + W_{21}x_{21} + W_{22}x_{22}$$




CNNs




CNNs




CNNs




CNNs




CNNs




CNNs




CNNs




CNNs





CNNs



"feature map"



CNNs




Note that we used the same weight matrix for the entire operation:

The idea that nearby input values are related is **built-in** to CNNs via “weight tying”.

CNNs

How do you “take the derivative” of that?



$$z_{ij} = \sum_{kl} W_{kl} x_{i+k, j+l}$$

CNNs

$$z_{ij} = \sum_{kl} W_{kl} x_{i+k, j+l}$$

Note: $\frac{\partial W_{ij}}{\partial W_{mn}} = \delta_{im} \delta_{jn}$

(weights are independent variables)

Let's find the values used to update weights W_{ij} :

$$\begin{aligned} dW_{mn} &= \left(\frac{\partial L}{\partial W_{mn}} \right) \epsilon = \sum_{ij} dz_{ij} \frac{\partial z_{ij}}{\partial W_{nm}} \\ &= \sum_{ij} dz_{ij} \frac{\partial W_{kl}}{\partial W_{nm}} x_{i+k, j+l} \\ &= \sum_{ij} dz_{ij} \delta_{kn} \delta_{lm} x_{i+k, j+l} \\ &= \sum_{ij} dz_{ij} x_{i+n, j+m} \\ &= \sum_{ij} dz_{ij} x_{n+i, m+j} = (x * dz)_{nm} \end{aligned}$$

CNNs

$$z_{ij} = \sum_{kl} W_{kl} x_{i+k, j+l}$$

Let's find the values used to update weights W_{ij} :



$$\begin{aligned} dW_{mn} &= \left(\frac{\partial L}{\partial W_{mn}} \right) \epsilon = \sum_{ij} dz_{ij} \frac{\partial z_{ij}}{\partial W_{nm}} \\ &= \sum_{ij} dz_{ij} \frac{\partial W_{kl}}{\partial W_{nm}} x_{i+k, j+l} \\ &= \sum_{ij} dz_{ij} \delta_{kn} \delta_{lm} x_{i+k, j+l} \\ &= \sum_{ij} dz_{ij} x_{i+n, j+m} \\ &= \sum_{ij} dz_{ij} x_{n+i, m+j} \end{aligned}$$

$$= (x * dz)_{nm}$$

The gradients for the weights in a convolution can be found by convolving the output gradients with the input.


CNNs

Some additional terminology:



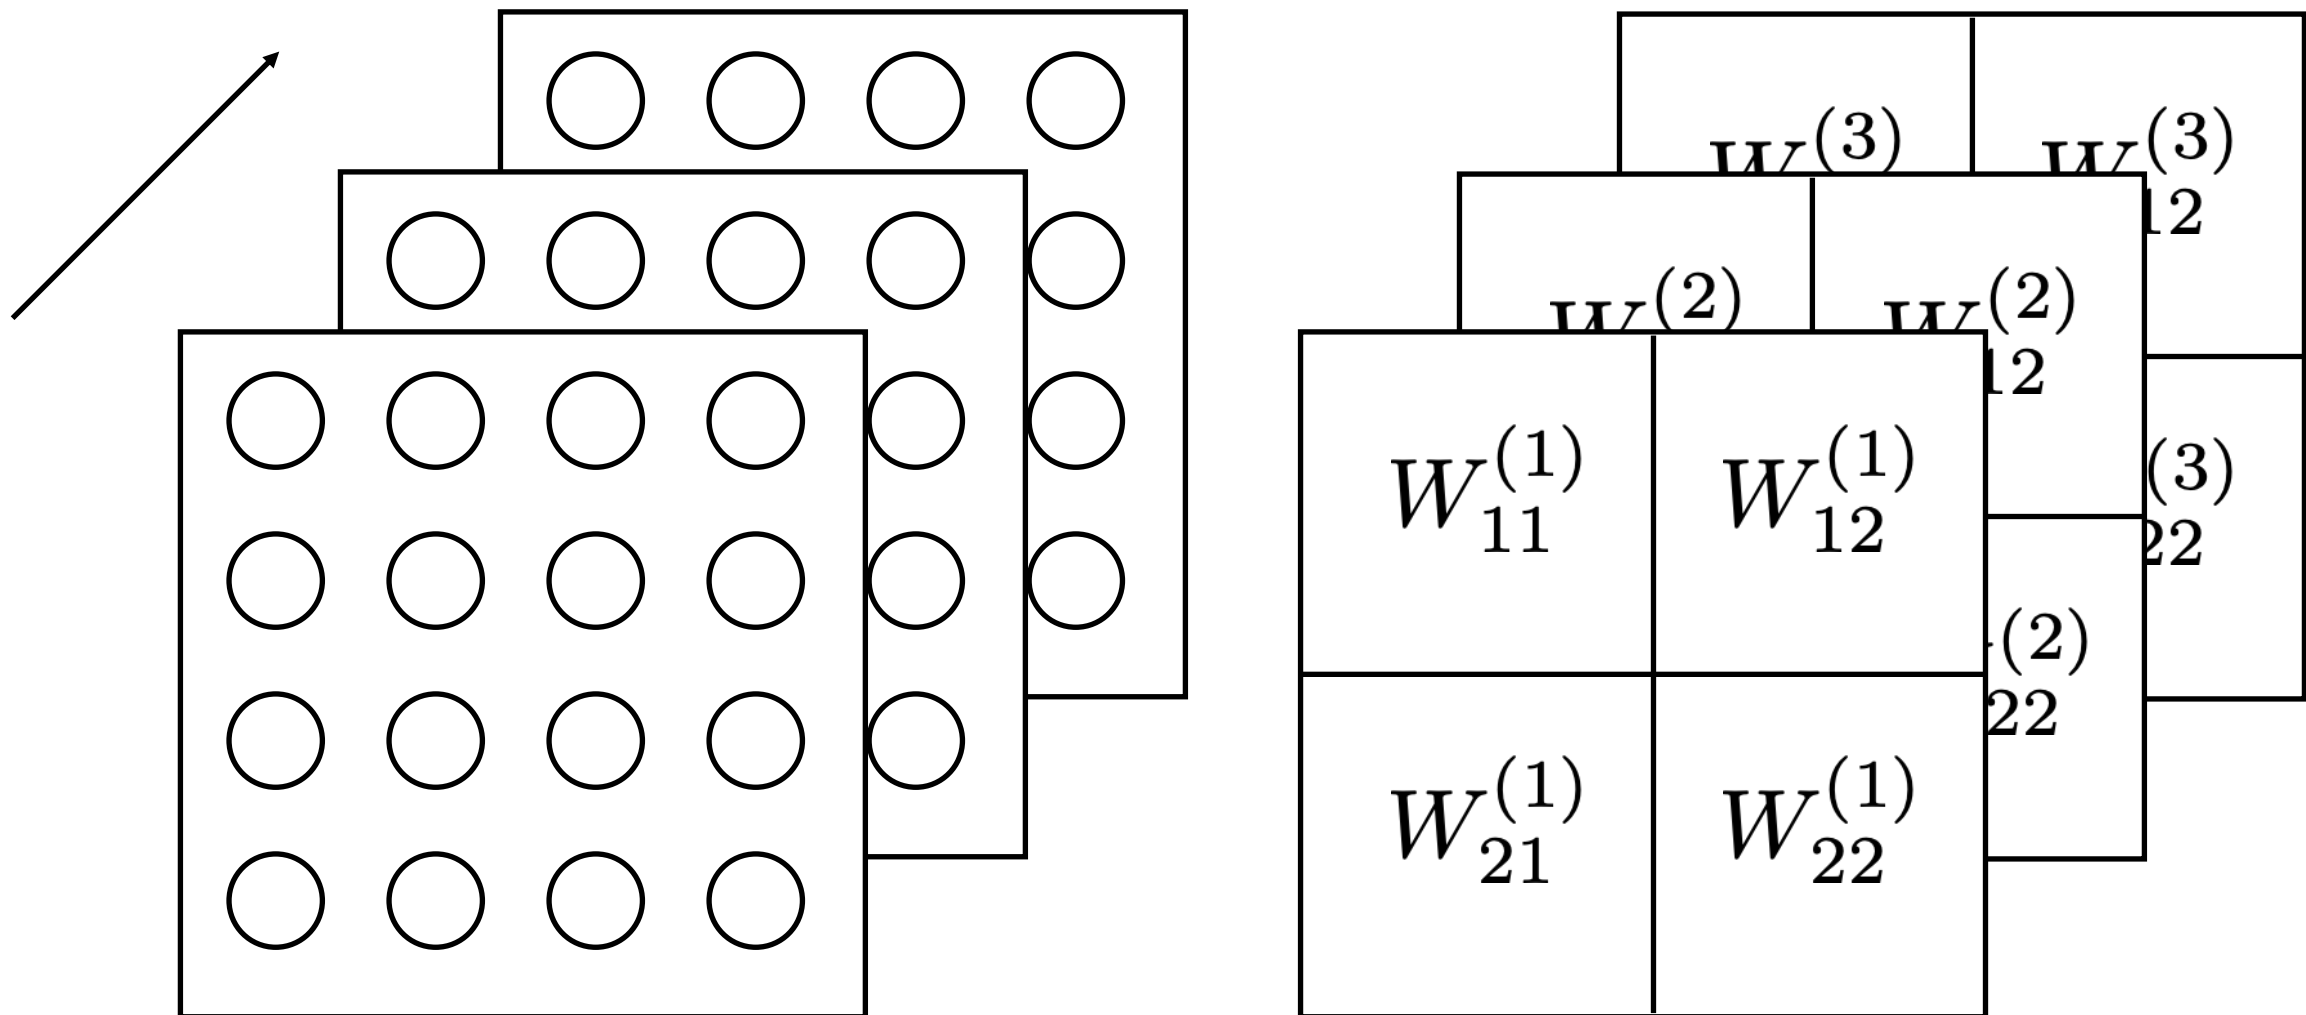
CNNs

(max/average) pooling



Selects only the maximum value, or
averages all values over a $K \times L$ interval


CNNs



Inputs and kernels can also have multiple *channels* (for example, 3 channels in an RGB image)

NEXT

Neutrino Experiment with a Xenon TPC




NEXT at Canfranc

- search for neutrinoless double-beta decay ($0\nu\beta\beta$)
- to be commissioned in 2020: 100 kg Xe, enriched to ^{136}Xe (90%)
- high pressure gas, electroluminescent TPC: capable of measuring **energy** of an interaction and reconstructing particle **tracks**

What are we looking for?

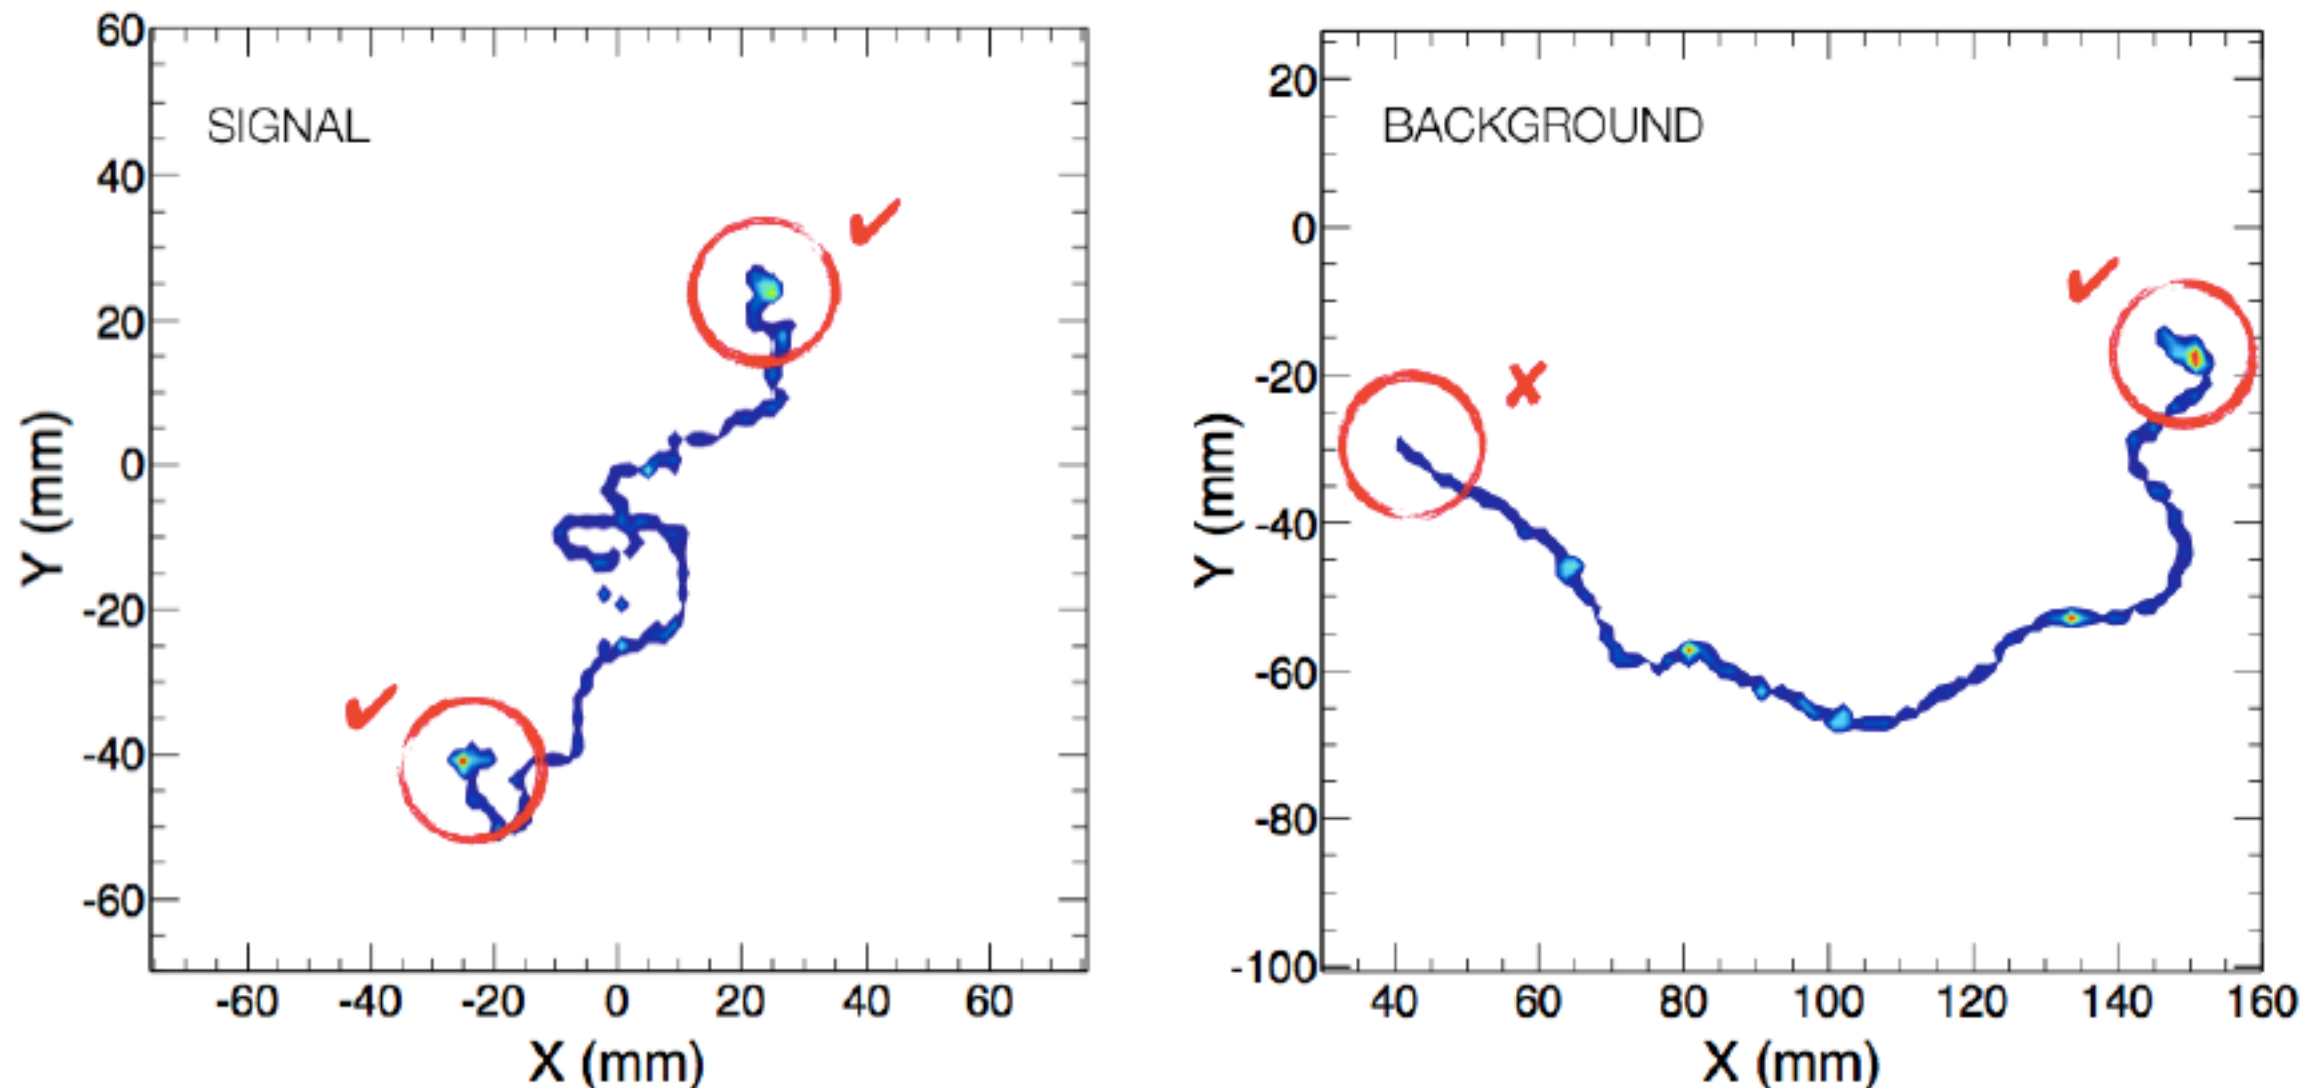
$0\nu\beta\beta$ would yield 2 electrons of total energy = $Q_{\beta\beta}$



- But we could still get background events that are not $0\nu\beta\beta$ but still fall into the energy peak

Topological signature in NEXT

JHEP 01, 104 (2016) [arXiv:1507.05902]




- Energetic electron leaves a high-density deposition at the end of its track (Bragg peak)
- Results in distinct topological signatures for **signal** and **background** events of the same energy

CNNs in Pytorch


NEXT simulation dataset (20x20x60 particle tracks):

- **Note these are actually e+e- pair events (similar topology and can be produced experimentally)**



CNNs in Pytorch

NEXT simulation dataset (20x20x60 particle tracks):



~ 35000 total events available/class, to be split up into training, validation, and test sets

We'll be dealing with tensors of shape:

$[N, 20, 20, 60]$

“batch size”

Note: Batches

Gradient descent steps are commonly performed with smaller subsets of the full training set called *minibatches* (or “batches”).

```
# Create a new Dataset
```

```
dataset_train = NEXTDataset(datafile_signal,  
datafile_background, nstart_train, nend_train)
```

```
# Create a new DataLoader
```


```
train_loader = DataLoader(dataset_train,  
batch_size=batch_size, shuffle=True)
```

At each step, the loss and gradients will be computed using `batch_size` data samples:

```
for batch_idx, (data, target) in enumerate(train_loader):  
  
    optimizer.zero_grad()  
    outputs = model(data)  
    loss = criterion(outputs, target)  
    loss.backward()  
    optimizer.step()
```


Note: Dropout

During training/backpropagation, some specified probability of removing all connections to a neuron for a step:



Note: Dropout

During training/backpropagation, some specified probability of removing all connections to a neuron for a step:



CNNs in Pytorch

Some relevant layers/commands

(see also <https://pytorch.org/docs/stable/nn.html>):

```
# Create a new linear ( $Wx + b$ ) layer  
nn.Linear(in_features, out_features)
```

```
# Create a new 2D convolutional layer  
torch.nn.Conv2d(C_in, C_out, kernel_size, stride, padding)
```

```
# Create a new 2D max pooling layer  
torch.nn.MaxPool2d(kernel_size, stride)
```

```
# Dropout layer with the specified probability  
torch.nn.Dropout(prob)
```

```
# Apply the sigmoid function  
torch.sigmoid(x)
```

```
# Apply the ReLU function  
torch.relu(x)
```

```
# Flatten a tensor (starting after the batch dimension)  
torch.flatten(x, start_dim=1)
```


CNNs in Pytorch

Some relevant layers/commands

(see also <https://pytorch.org/docs/stable/nn.html>):

```
# Create a new linear (Wx + b) layer  
nn.Linear(in_features, out_features)
```

```
# Create a new 2D convolutional layer  
torch.nn.Conv2d(C_in, C_out, kernel_size, stride, padding)
```

```
# Create a new 2D max pooling layer  
torch.nn.MaxPool2d(kernel_size, stride)
```

```
# Dropout layer with the specified probability  
torch.nn.Dropout(prob)
```

```
# Apply the sigmoid function  
torch.sigmoid(x)
```

```
# Apply the ReLU function  
torch.relu(x)
```

```
# Flatten a tensor (starting after the batch dimension)  
torch.flatten(x, start_dim=1)
```

ReLU (rectified linear units) is another common activation used in CNNs

$\text{relu}(x) = \max(0, x)$

CNNs in Pytorch

An example net (not a CNN):

```
class FCNet(nn.Module):
```

```
    def __init__(self):  
        super(FCNet, self).__init__()  
        self.fc1 = nn.Linear(xdim*ydim*zdim, 32)  
        self.fc2 = nn.Linear(32, 1)
```

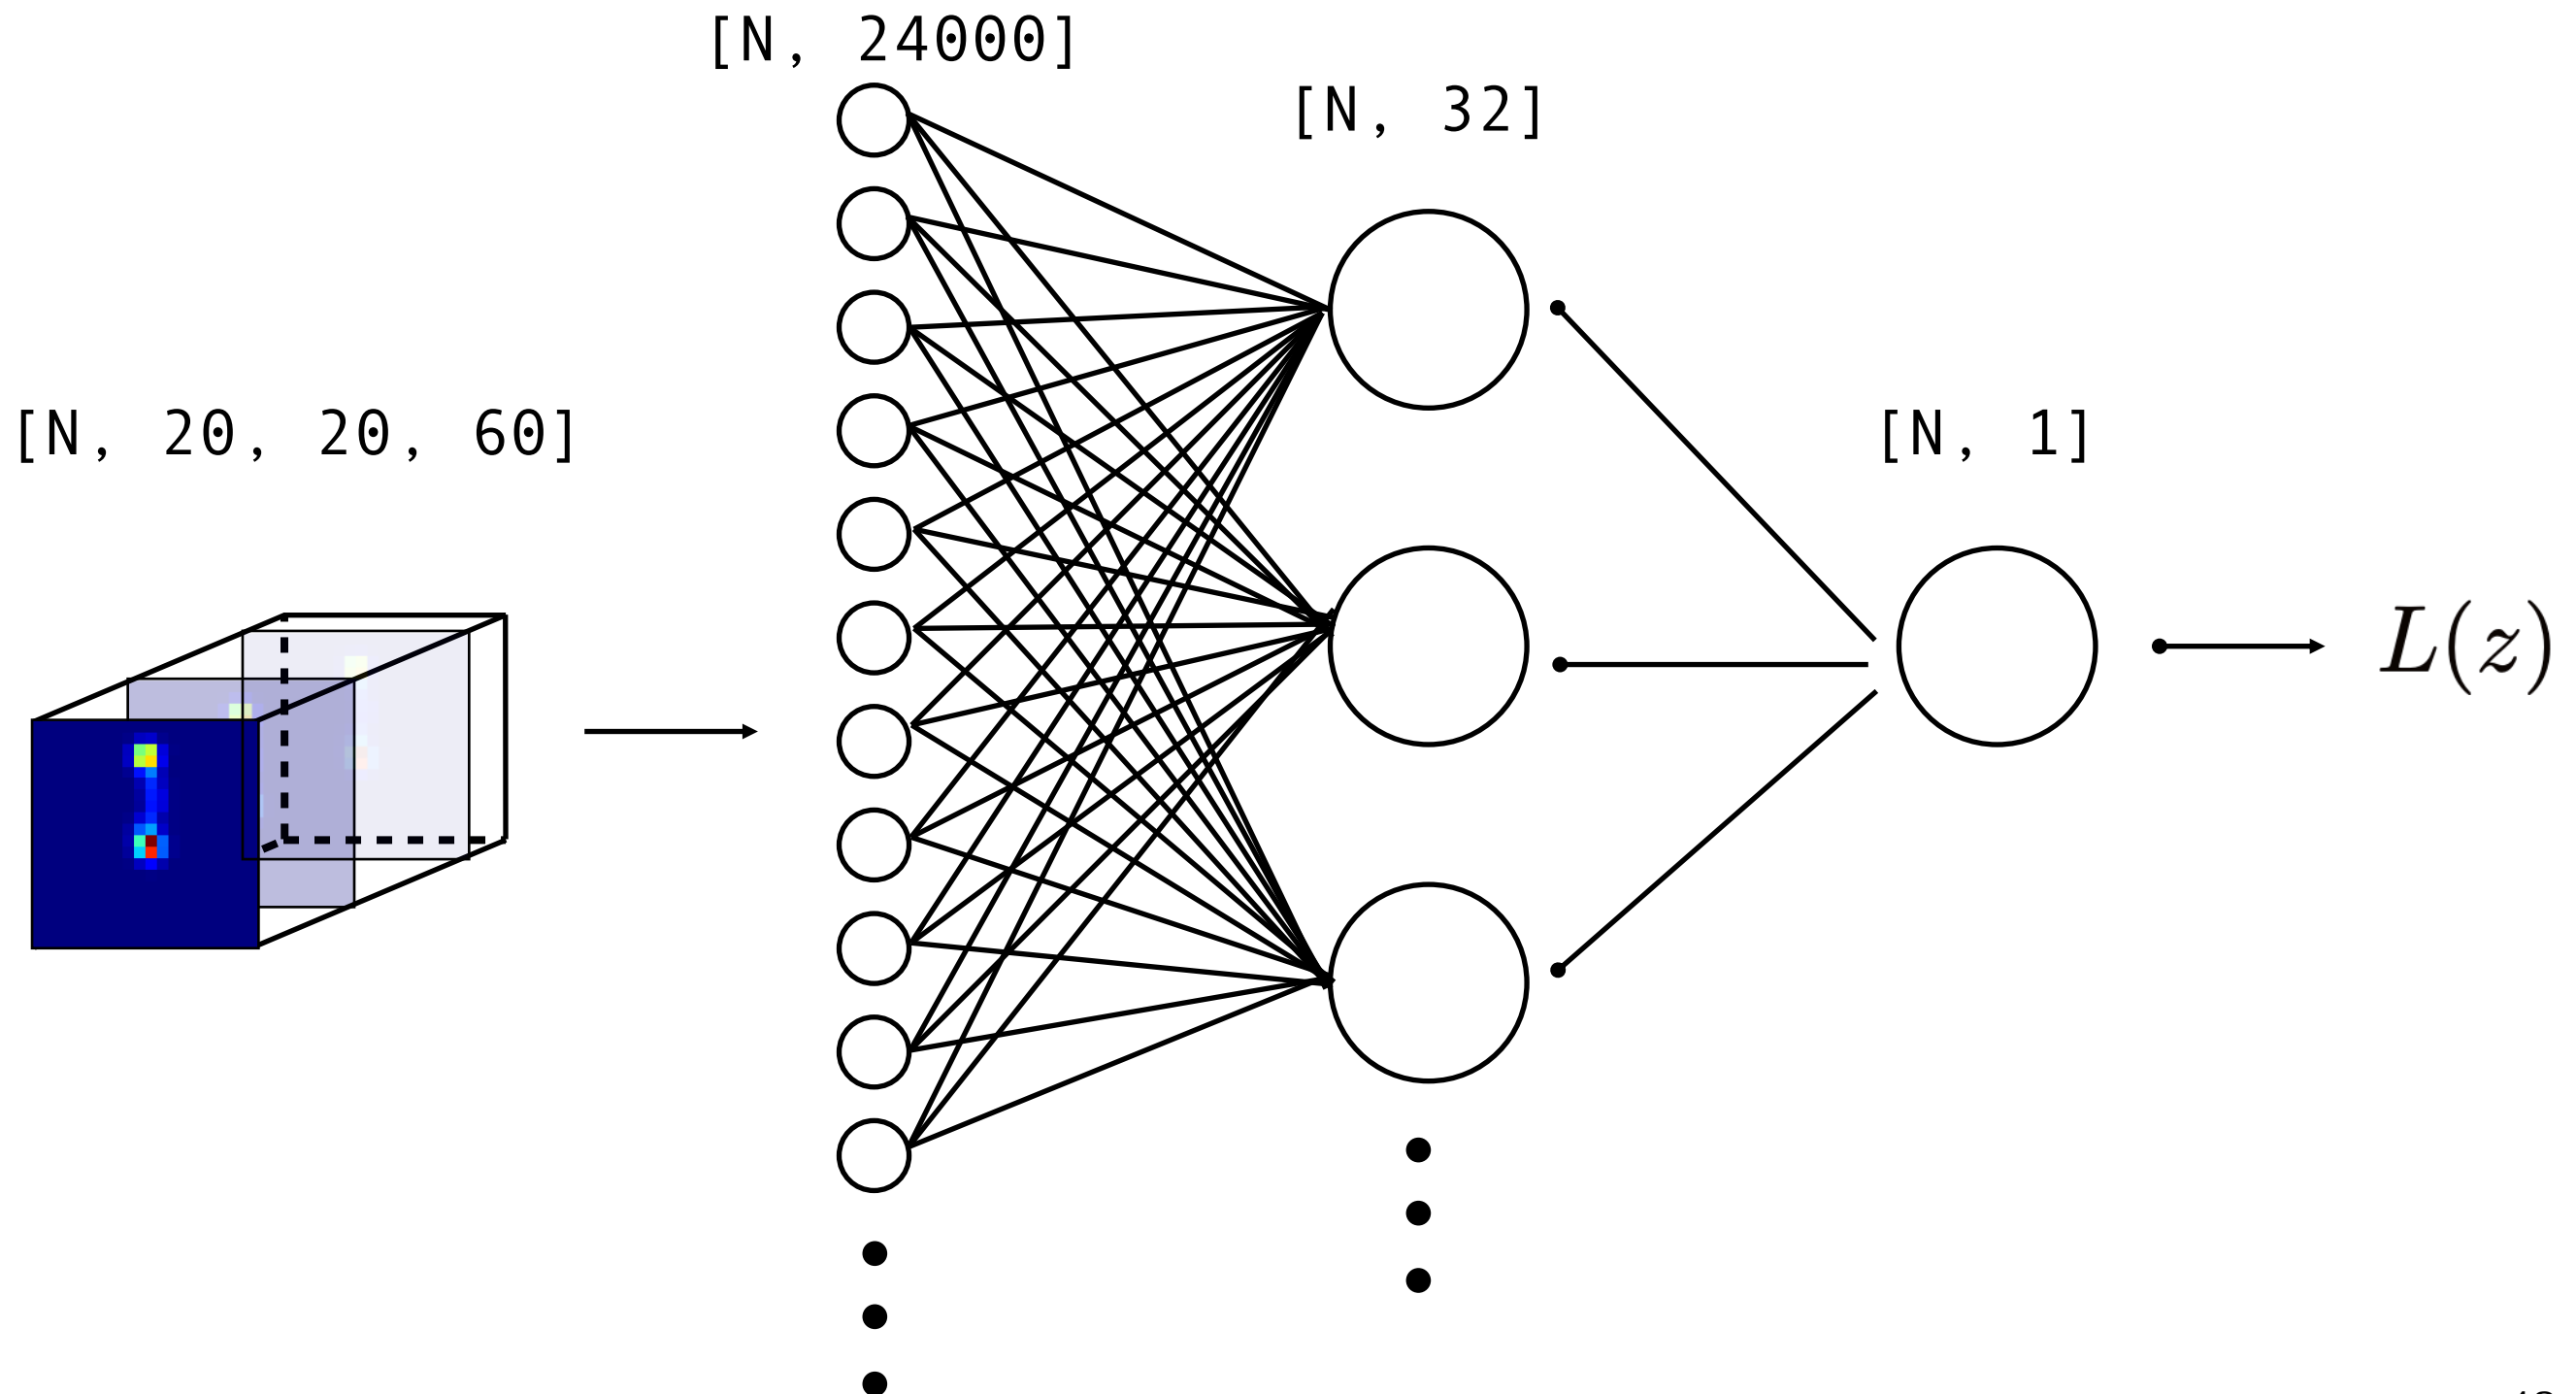
```
    def forward(self, x):  
        x = torch.flatten(x, start_dim=1)  
        x = self.fc1(x)  
        x = torch.sigmoid(x)  
        x = self.fc2(x)  
        return x
```

define layers

apply layers



CNNs in Pytorch

An example net (not convolutional):



CNNs in Pytorch

Performance: **can this be improved with a CNN?**



https://colab.research.google.com/github/jerenner/uscnncourse/blob/master/next/NEXT_classification.ipynb

EXTRA SLIDES

Neural networks

Backpropagation on a graph is essentially the chain rule:

$$\begin{aligned}\frac{\partial L}{\partial W_i} &= \sum_j \frac{\partial L}{\partial z_j} \sum_k \frac{\partial z_j}{\partial y_k} \frac{\partial y_k}{\partial W_i} = \sum_{j,k} \frac{\partial L}{\partial z_j} \frac{\partial z_j}{\partial y_k} \frac{\partial y_k}{\partial W_i} = \sum_{j,k} \frac{\partial y_k}{\partial W_i} \frac{\partial z_j}{\partial y_k} \frac{\partial L}{\partial z_j} \\&= \sum_{j,k} \frac{\partial y_k}{\partial W_i} \left[\frac{\partial z_j}{\partial y_k} \left(\frac{\partial L}{\partial z_j} \right) \right] = \sum_k \frac{\partial y_k}{\partial W_i} \sum_j \left[\frac{\partial z_j}{\partial y_k} dz_j \right], \text{ where } dz_j = \frac{\partial L}{\partial z_j} \\&= \sum_k \frac{\partial y_k}{\partial W_i} dy_k, \text{ where } dy_k = \sum_j \left[\frac{\partial z_j}{\partial y_k} dz_j \right]\end{aligned}$$