Convolutional Neural Networks and Applications in the NEXT Experiment

G. Díaz, J.A. Hernando, M. Kekic, J. Renner

IFGAE / Universidade de Santiago de Compostela
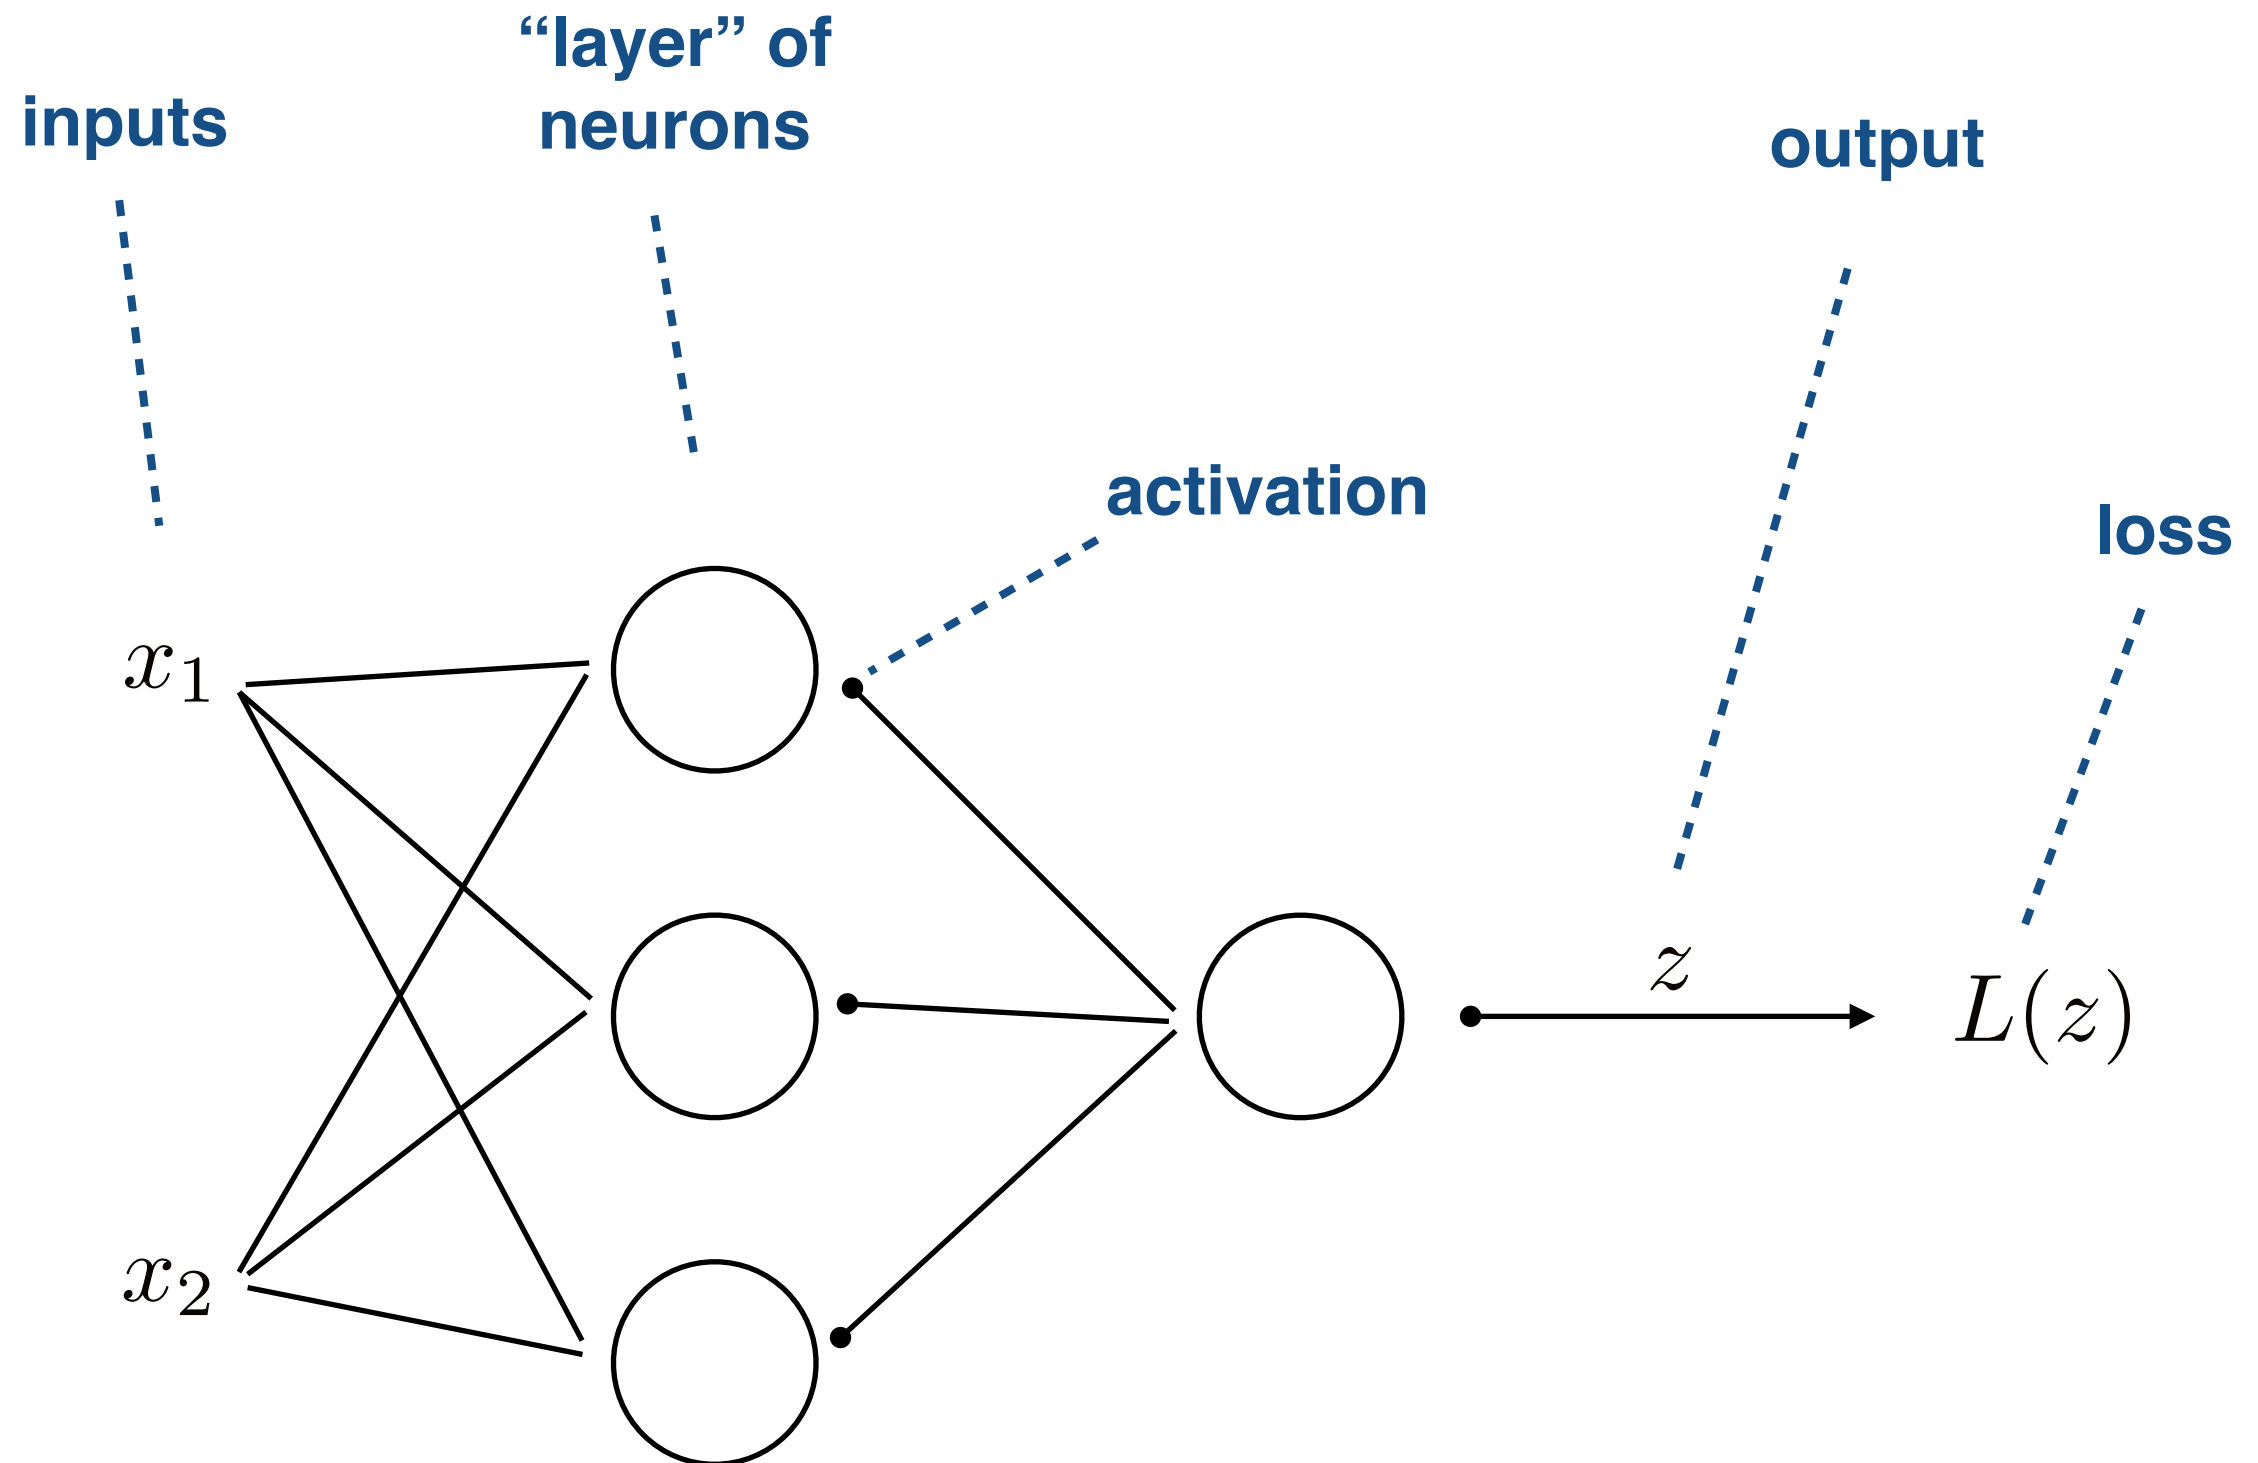
Advanced Computing and Machine Learning

June 1, 2020

Review: last week

Neural networks

Neural networks

inputs

"layer" of neurons

output

activation

loss

x_1

x_2

z

$L(z)$

4

Neural networks

Each neuron contains a **weight** corresponding to each input, and a **bias**:

$$W_1 x_1 + W_2 x_2 + b$$

y

x_1

x_2

Activation function applied to introduce nonlinearity:

$$y = \sigma(W_1 x_1 + W_2 x_2 + b) = \frac{1}{1 + e^{-(W_1 x_1 + W_2 x_2 + b)}}$$

Neural networks

To minimize L, we want to push the entire network (all parameters) in the direction of -gradient(L):

$$W_1 \rightarrow W_1 - \left(\frac{\partial L}{\partial W_1} \right) \epsilon$$

z

$L(z)$

Backpropagate to get $\left(\dfrac{\partial L}{\partial W_1} \right)$

Neural networks

$$\left(\frac{\partial L}{\partial W_1}\right)\epsilon = \left(\frac{\partial y}{\partial W_1}\right)\left(\frac{\partial z}{\partial y}\right)\left(\frac{\partial L}{\partial z}\right)\epsilon$$

`loss.backward()`

x_1

$\left(\frac{\partial z}{\partial y}\right)\left(\frac{\partial L}{\partial z}\right)\epsilon$

y

z

$L(z)$

$\left(\frac{\partial L}{\partial z}\right)\epsilon$

ϵ

x_2

Note: ∂z/∂y will have to account for the functional forms of: **this neuron**

and this activation

Neural networks

- Note that backpropagation on graphs can be performed on a node-by-node basis:

Neural networks

Neural networks

$$x_1 \qquad y_1$$
$$dx_1 \qquad dy_1$$
$$f(\vec{x}; W)$$
$$dx_2 \qquad dy_2$$
$$x_2 \qquad y_2$$

- Given gradients dy$_1$ and dy$_2$
- Given f differentiable
- Compute gradients dW for parameter updates, and gradients dx$_1$ and dx$_2$ to continue backpropagation

$$dx_i = \sum_k \frac{\partial y_k}{\partial x_i} dy_k$$

Libraries such as PyTorch handle the mechanics of this automatically

Neural networks

- So far, we've considered:

$$f(\vec{x}, W) = W\vec{x} + \vec{b}$$

Let's try something else!

Convolutional Neural Networks (CNNs)

CNNs

Inputs, arranged in
a 2D matrix

W_{11} W_{12}

W_{21} W_{22}

KxL kernel of weights
(in this case 2x2)

CNNs

Multiply and add:

$$z_{11} = W_{11}x_{11} + W_{12}x_{12} + W_{21}x_{21} + W_{22}x_{22}$$

CNNs

CNNs

CNNs

CNNs

CNNs

CNNs

CNNs

CNNs

CNNs

"feature map"

CNNs

$$W_{11} \quad W_{12}$$
$$W_{21} \quad W_{22}$$

Note that we used the same weight matrix for the entire operation:

The idea that nearby input values are related is **built-in** to CNNs via "weight tying".

CNNs

How do you "take the derivative" of that?

$$W$$

$$x$$

$$z = x * W$$

$$*$$

$$dx$$

$$dz$$

$$z_{ij} = \sum_{kl} W_{kl} x_{i+k,j+l}$$

CNNs

$$z_{ij} = \sum_{kl} W_{kl} x_{i+k,j+l}$$

Note: $\dfrac{\partial W_{ij}}{\partial W_{mn}} = \delta_{im}\delta_{jn}$

(weights are independent variables)

Let's find the values used to update weights W_{ij}:

$$dW_{mn} = \left(\frac{\partial L}{\partial W_{mn}}\right)\epsilon = \sum_{ij} dz_{ij} \frac{\partial z_{ij}}{\partial W_{nm}}$$

$$= \sum_{ij} dz_{ij} \frac{\partial W_{kl}}{\partial W_{nm}} x_{i+k,j+l}$$

$$= \sum_{ij} dz_{ij} \delta_{kn}\delta_{lm} x_{i+k,j+l}$$

$$= \sum_{ij} dz_{ij} x_{i+n,j+m}$$

$$= \sum_{ij} dz_{ij} x_{n+i,m+j} \boxed{= (x * dz)_{nm}}$$

CNNs

$$z_{ij} = \sum_{kl} W_{kl} x_{i+k,j+l}$$

Let's find the values used to update weights W_{ij}:

$$dW_{mn} = \left(\frac{\partial L}{\partial W_{mn}} \right) \epsilon = \sum_{ij} dz_{ij} \frac{\partial z_{ij}}{\partial W_{nm}}$$

$$= \sum_{ij} dz_{ij} \frac{\partial W_{kl}}{\partial W_{nm}} x_{i+k,j+l}$$

$$= \sum_{ij} dz_{ij} \delta_{kn} \delta_{lm} x_{i+k,j+l}$$

$$= \sum_{ij} dz_{ij} x_{i+n,j+m}$$

$$= \sum_{ij} dz_{ij} x_{n+i,m+j} \boxed{= (x * dz)_{nm}}$$

The gradients for the weights in a convolution can be found by convolving the output gradients with the input.

CNNs

Some additional terminology:

stride

padding (0s)

CNNs

(max/average) pooling

Selects only the maximum value, or
averages all values over a KxL interval

CNNs
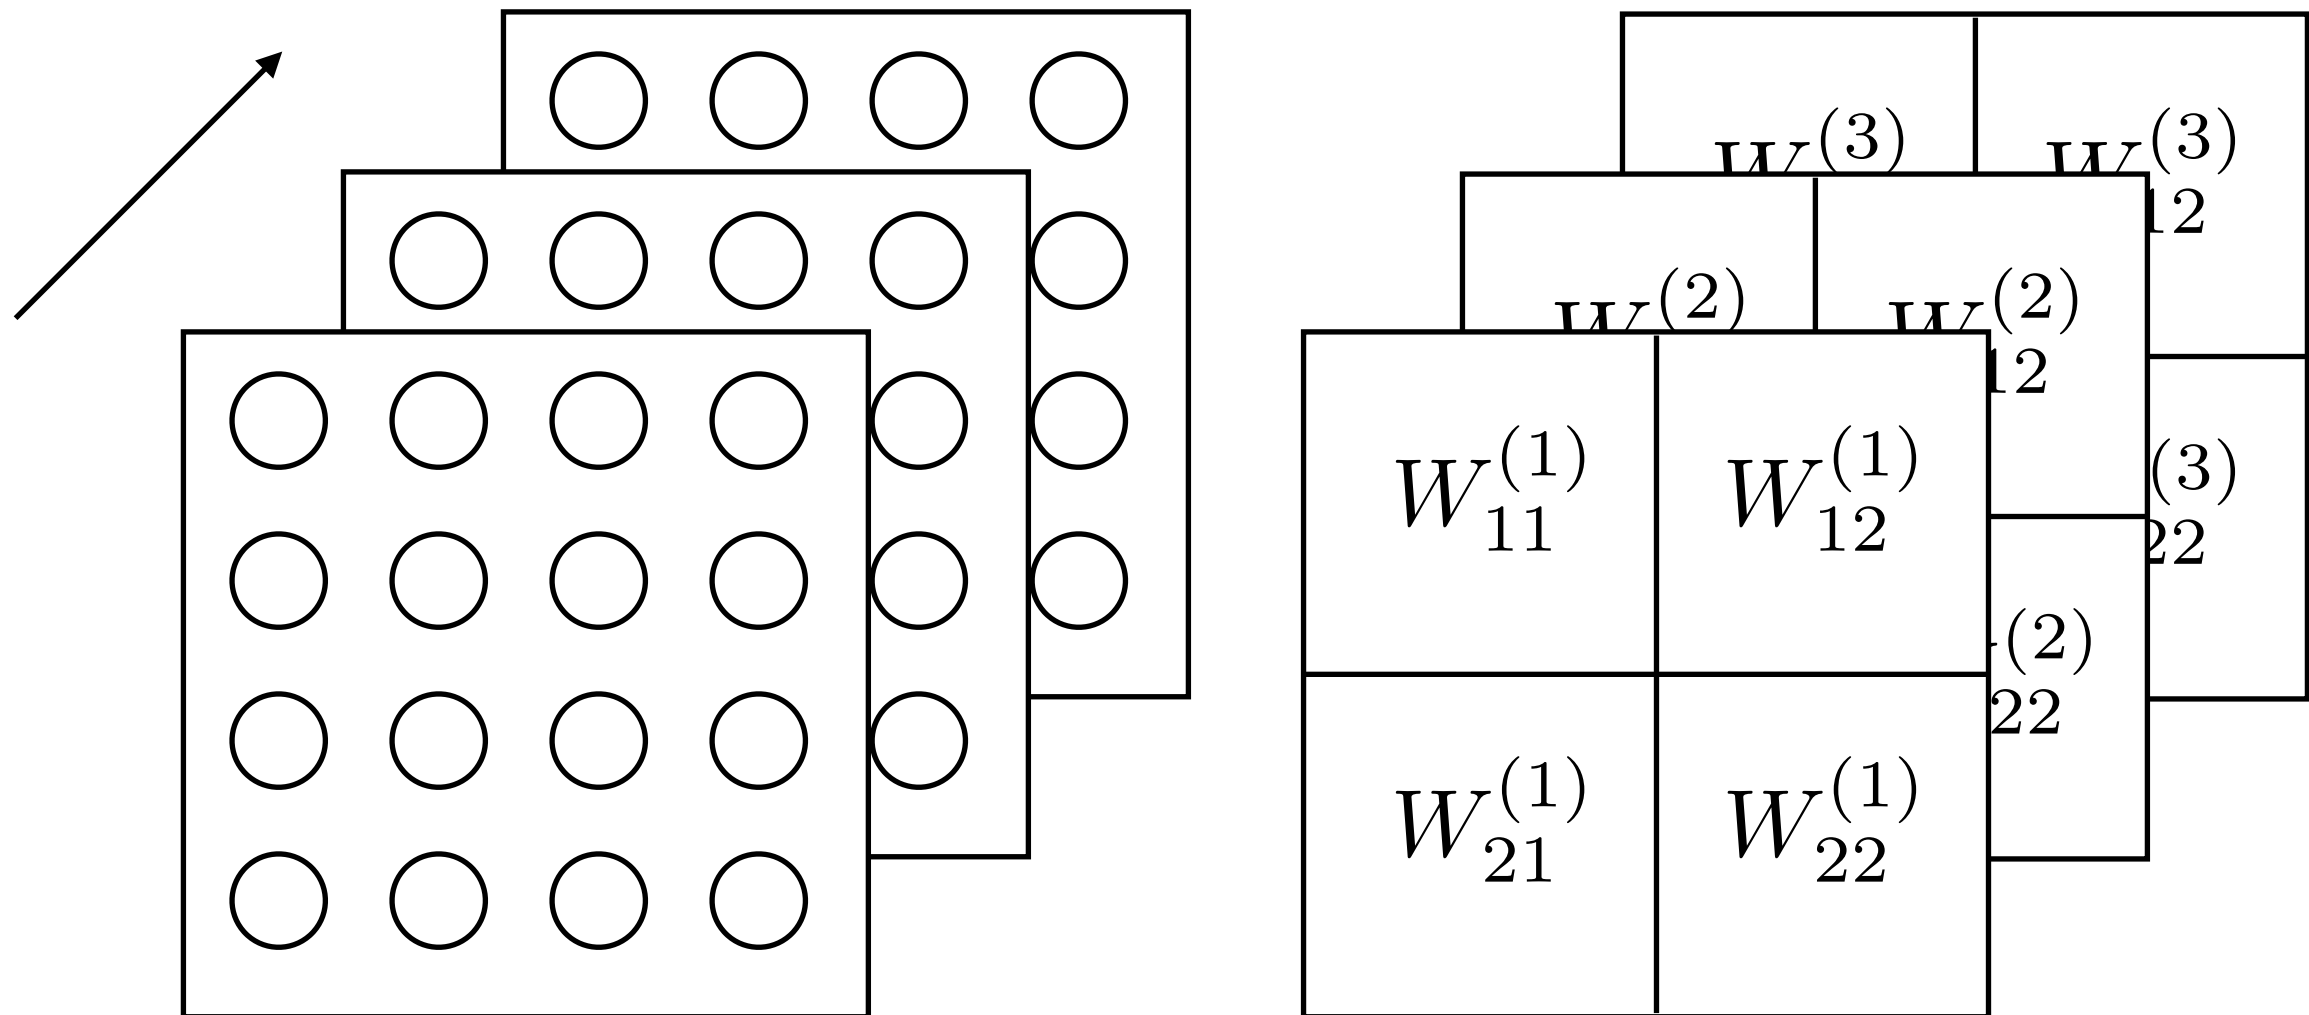
Inputs and kernels can also have multiple *channels* (for example, 3 channels in an RGB image)

NEXT

Neutrino
Experiment with a
Xenon
TPC

NEXT at Canfranc

- search for neutrinoless double-beta decay (0νßß)

- to be commissioned in 2020: 100 kg Xe, enriched to ^{136}Xe (90%)

- high pressure gas, electroluminescent TPC: capable of measuring **energy** of an interaction and reconstructing particle **tracks**

What are we looking for?

0vßß would yield 2 electrons of total energy = Q$_{ßß}$

- But we could still get background events that are not 0vßß but still fall into the energy peak

Topological signature in NEXT

- Energetic electron leaves a high-density deposition at the end of its track (Bragg peak)

- Results in distinct topological signatures for **signal** and **background** events of the same energy

CNNs in Pytorch

NEXT simulation dataset (20x20x60 particle tracks):

- **Note these are actually e+e- pair events (similar topology and can be produced experimentally)**

CNNs in Pytorch

NEXT simulation dataset (20x20x60 particle tracks):

~ 35000 total events available/class, to be split up into training, validation, and test sets

We'll be dealing with tensors of shape:

`[N, 20, 20, 60]`

"batch size"

Note: Batches

Gradient descent steps are commonly performed with smaller subsets of the full training set called *minibatches* (or "batches").

```
# Create a new Dataset
dataset_train = NEXTDataset(datafile_signal,
datafile_background, nstart_train, nend_train)

# Create a new DataLoader
train_loader  = DataLoader(dataset_train,
batch_size=batch_size, shuffle=True)
```

At each step, the loss and gradients will be computed using `batch_size` data samples:

```
for batch_idx, (data, target) in enumerate(train_loader):

  optimizer.zero_grad()
  outputs = model(data)
  loss = criterion(outputs, target)
  loss.backward()
  optimizer.step()
```

Note: Dropout

During training/backpropagation, some specified probability
of removing all connections to a neuron for a step:

$L(z)$

Note: Dropout

During training/backpropagation, some specified probability
of removing all connections to a neuron for a step:

x_1

x_2

$L(z)$

— Helps avoid overfitting

CNNs in Pytorch

Some relevant layers/commands
(see also https://pytorch.org/docs/stable/nn.html):

```python
# Create a new linear (Wx + b) layer
nn.Linear(in_features, out_features)

# Create a new 2D convolutional layer
torch.nn.Conv2d(C_in, C_out, kernel_size, stride, padding)

# Create a new 2D max pooling layer
torch.nn.MaxPool2d(kernel_size, stride)

# Dropout layer with the specified probability
torch.nn.Dropout(prob)


# Apply the sigmoid function
torch.sigmoid(x)

# Apply the ReLU function
torch.relu(x)

# Flatten a tensor (starting after the batch dimension)
torch.flatten(x, start_dim=1)
```

CNNs in Pytorch

Some relevant layers/commands
(see also https://pytorch.org/docs/stable/nn.html):

```python
# Create a new linear (Wx + b) layer
nn.Linear(in_features, out_features)

# Create a new 2D convolutional layer
torch.nn.Conv2d(C_in, C_out, kernel_size, stride, padding)

# Create a new 2D max pooling layer
torch.nn.MaxPool2d(kernel_size, stride)

# Dropout layer with the specified probability
torch.nn.Dropout(prob)


# Apply the sigmoid function
torch.sigmoid(x)

# Apply the ReLU function
torch.relu(x)

# Flatten a tensor (starting after the batch dimension)
torch.flatten(x, start_dim=1)
```

ReLU (rectified linear units) is another common activation used in CNNs

relu(x) = max(0,x)

CNNs in Pytorch

An example net (not a CNN):

```python
class FCNet(nn.Module):

    def __init__(self):
        super(FCNet, self).__init__()
        self.fc1 = nn.Linear(xdim*ydim*zdim, 32)
        self.fc2 = nn.Linear(32, 1)

    def forward(self, x):
        x = torch.flatten(x, start_dim=1)
        x = self.fc1(x)
        x = torch.sigmoid(x)
        x = self.fc2(x)
        return x
```
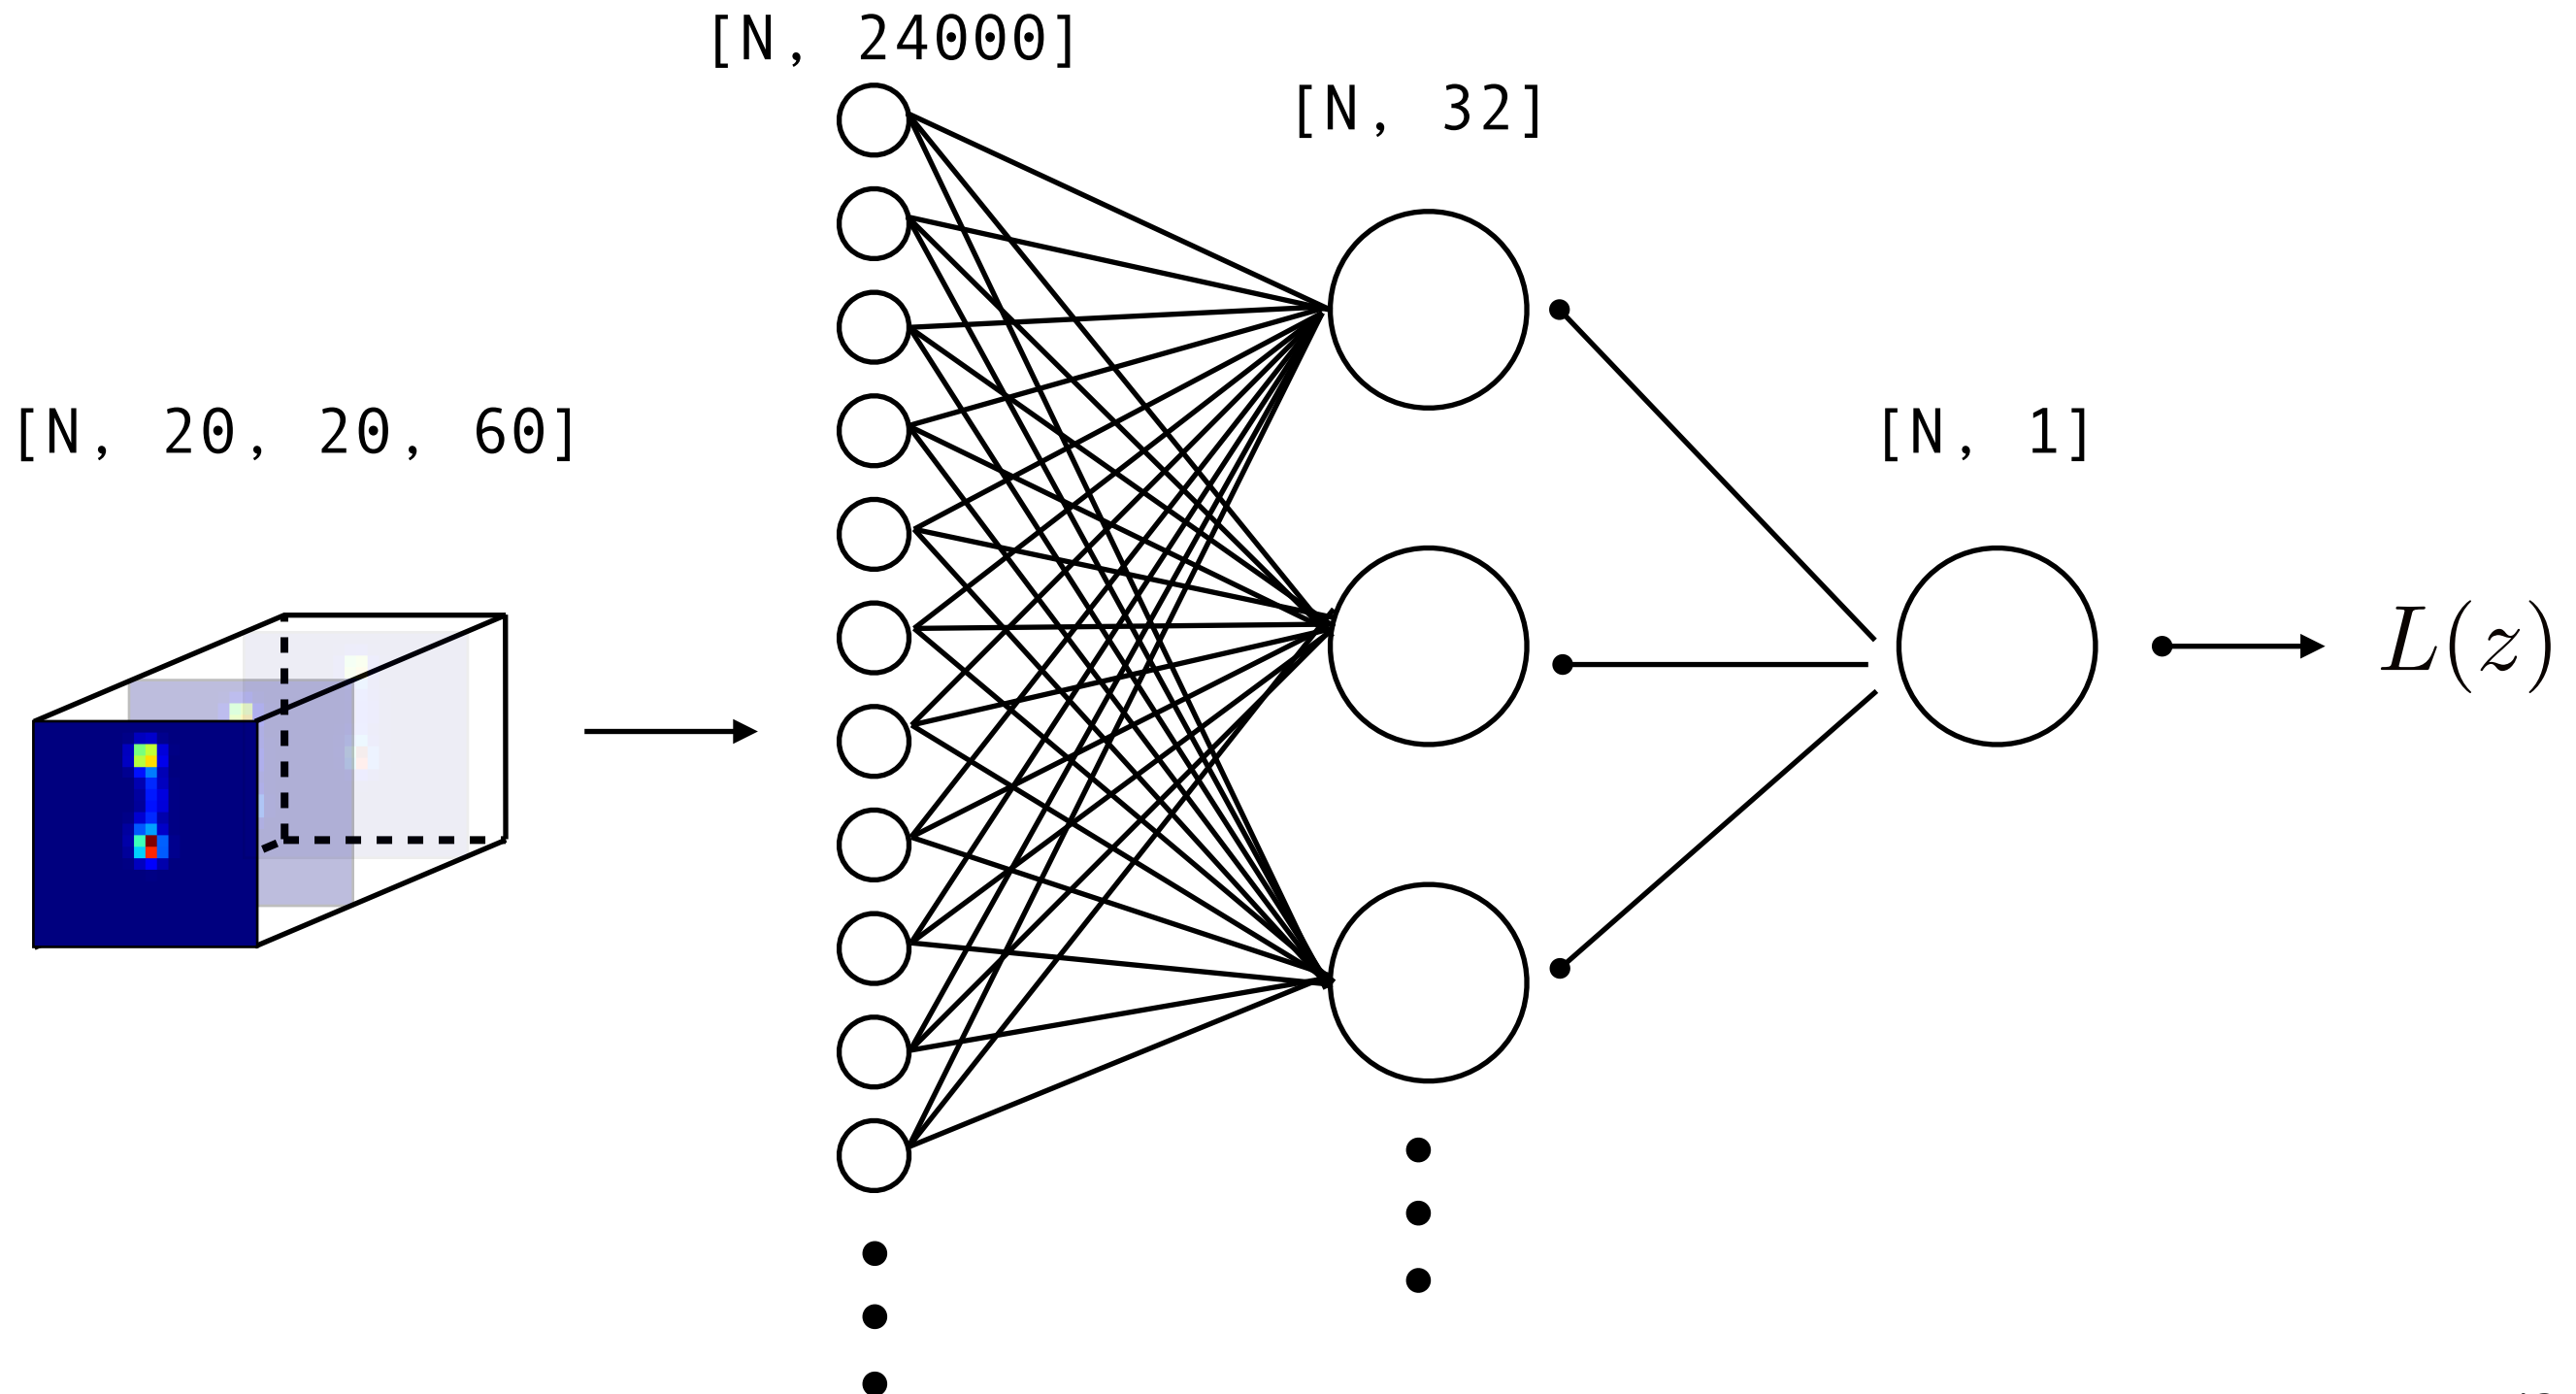
define layers

apply layers

CNNs in Pytorch

An example net (not convolutional):

[N, 24000]

[N, 32]

[N, 20, 20, 60]

[N, 1]

$L(z)$

CNNs in Pytorch

Performance: **can this be improved with a CNN?**

**https://colab.research.google.com/github/jerenner/
uscnncourse/blob/master/next/NEXT_classification.ipynb**

EXTRA SLIDES

Neural networks

Backpropagation on a graph is essentially the chain rule:

$$\frac{\partial L}{\partial W_i} = \sum_j \frac{\partial L}{\partial z_j} \sum_k \frac{\partial z_j}{\partial y_k} \frac{\partial y_k}{\partial W_i} = \sum_{j,k} \frac{\partial L}{\partial z_j} \frac{\partial z_j}{\partial y_k} \frac{\partial y_k}{\partial W_i} = \sum_{j,k} \frac{\partial y_k}{\partial W_i} \frac{\partial z_j}{\partial y_k} \frac{\partial L}{\partial z_j}$$

$$= \sum_{j,k} \frac{\partial y_k}{\partial W_i} \left[\frac{\partial z_j}{\partial y_k} \left(\frac{\partial L}{\partial z_j} \right) \right] = \sum_k \frac{\partial y_k}{\partial W_i} \sum_j \left[\frac{\partial z_j}{\partial y_k} dz_j \right], \text{where } dz_j = \frac{\partial L}{\partial z_j}$$

$$= \sum_k \frac{\partial y_k}{\partial W_i} dy_k, \text{where } dy_k = \sum_j \left[\frac{\partial z_j}{\partial y_k} dz_j \right]$$